

## [Question 1.1]

### Implementation:

```
first1.1.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn.model_selection import train_test_split, learning_curve
4  from sklearn.tree import DecisionTreeClassifier
5  from sklearn.linear_model import LogisticRegression
6  from sklearn.metrics import accuracy_score
7  import pandas as pd
8  from sklearn.datasets import load_iris
9
10 # Load the Iris dataset
11 data = load_iris()
12 X = data.data[:, :2] # Use the first two features for simplicity
13 y = data.target
14
15 # Split the dataset
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
17
18 # Save the dataset to a CSV file
19 dataset = pd.DataFrame(X, columns=['SepalLength', 'SepalWidth'])
20 dataset['Label'] = y
21 dataset.to_csv('iris_dataset.csv', index=False)
22
23 # Underfitting model: Logistic Regression
24 underfit_model = LogisticRegression()
25 underfit_model.fit(X_train, y_train)
26
27 # Overfitting model: Decision Tree with high depth
28 overfit_model = DecisionTreeClassifier(max_depth=10)
29 overfit_model.fit(X_train, y_train)
30
31 # Balanced model: Decision Tree with moderate depth
32 balanced_model = DecisionTreeClassifier(max_depth=3)
33 balanced_model.fit(X_train, y_train)
34
35 # Accuracy Scores
36 underfit_train_acc = accuracy_score(y_train, underfit_model.predict(X_train))
37 underfit_test_acc = accuracy_score(y_test, underfit_model.predict(X_test))
38
39 overfit_train_acc = accuracy_score(y_train, overfit_model.predict(X_train))
40 overfit_test_acc = accuracy_score(y_test, overfit_model.predict(X_test))
41
42 balanced_train_acc = accuracy_score(y_train, balanced_model.predict(X_train))
43 balanced_test_acc = accuracy_score(y_test, balanced_model.predict(X_test))
44
45 # Visualization function
46 def plot_decision_boundary(model, X, y, title):
47     x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
48     y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
49     xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01), np.arange(y_min, y_max, 0.01))
50     Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
51     Z = Z.reshape(xx.shape)
52     plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
53     plt.scatter(X[:, 0], X[:, 1], c=y, edgecolor='k', cmap=plt.cm.Paired)
54     plt.title(title)
55
```

```

55
56 # Plot decision boundaries
57 plt.figure(figsize=(15, 5))
58
59 plt.subplot(1, 3, 1)
60 plot_decision_boundary(underfit_model, X, y, "Underfitting: Logistic Regression")
61
62 plt.subplot(1, 3, 2)
63 plot_decision_boundary(overfit_model, X, y, "Overfitting: Decision Tree (High Depth)")
64
65 plt.subplot(1, 3, 3)
66 plot_decision_boundary(balanced_model, X, y, "Balanced Fit: Decision Tree (Moderate Depth)")
67
68 plt.tight_layout()
69 plt.show()
70
71 # Bias-Variance tradeoff analysis
72 print("Dataset saved as 'iris_dataset.csv'")
73 print("Underfitting Model: Train Accuracy =", underfit_train_acc, ", Test Accuracy =", underfit_test_acc)
74 print("Overfitting Model: Train Accuracy =", overfit_train_acc, ", Test Accuracy =", overfit_test_acc)
75 print("Balanced Model: Train Accuracy =", balanced_train_acc, ", Test Accuracy =", balanced_test_acc)
76
77 # Learning Curve for Bias-Variance Trade-off
78 def plot_learning_curve(model, title):
79     train_sizes, train_scores, test_scores = learning_curve(model, X, y, train_sizes=np.linspace(0.1, 1.0, 10), cv=5)
80
81     train_mean = np.mean(train_scores, axis=1)
82     test_mean = np.mean(test_scores, axis=1)
83
84     plt.plot(train_sizes, train_mean, label="Training Accuracy")
85     plt.plot(train_sizes, test_mean, label="Testing Accuracy")
86     plt.title(title)
87     plt.xlabel("Training Size")
88     plt.ylabel("Accuracy")
89     plt.legend()
90     plt.grid(True)
91
92 # Plot learning curves for all three models
93 plt.figure(figsize=(15, 10))
94
95 plt.subplot(2, 2, 1)
96 plot_learning_curve(underfit_model, "Learning Curve: Logistic Regression (Underfitting)")
97
98 plt.subplot(2, 2, 2)
99 plot_learning_curve(overfit_model, "Learning Curve: Decision Tree (Overfitting)")
100
101 plt.subplot(2, 2, 3)
102 plot_learning_curve(balanced_model, "Learning Curve: Decision Tree (Balanced)")
103
104 plt.tight_layout()
105 plt.show()
106

```

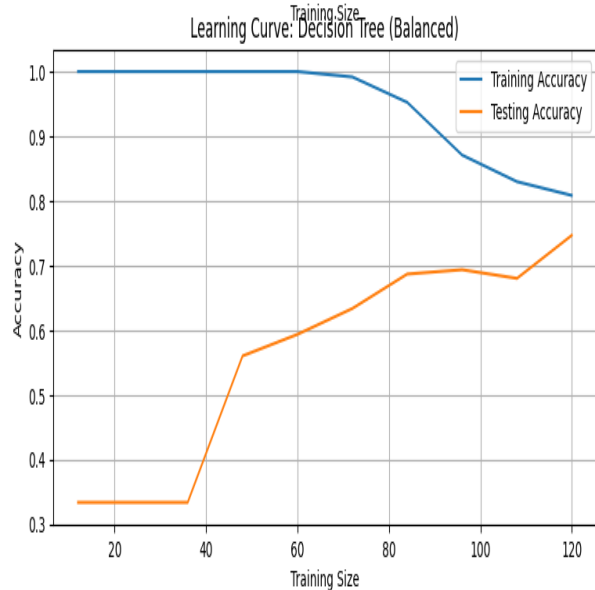
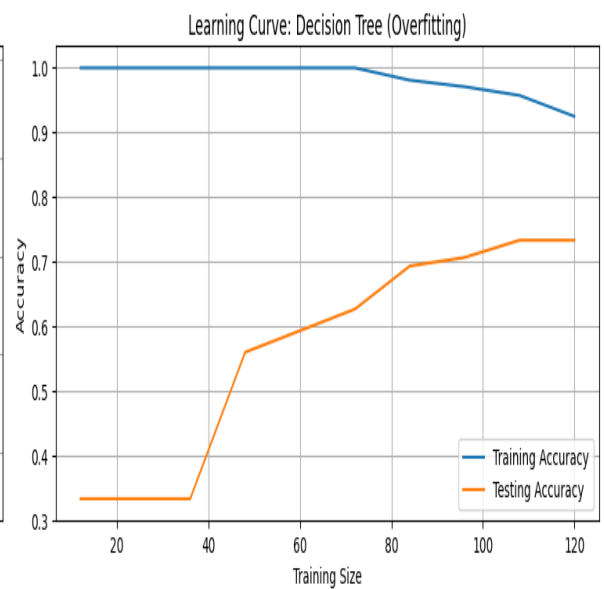
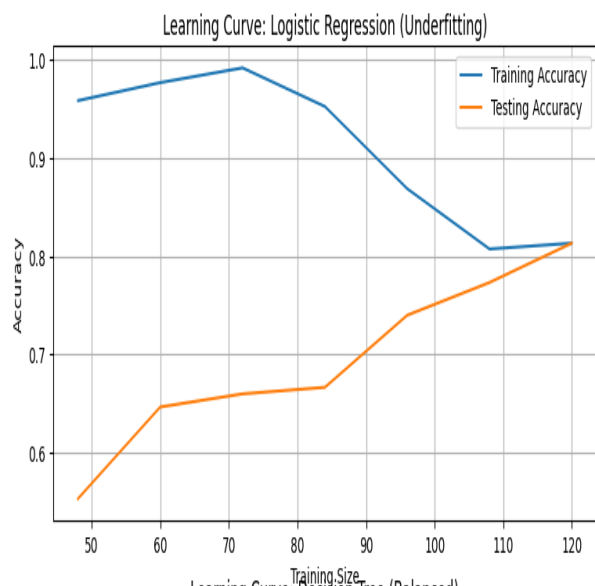
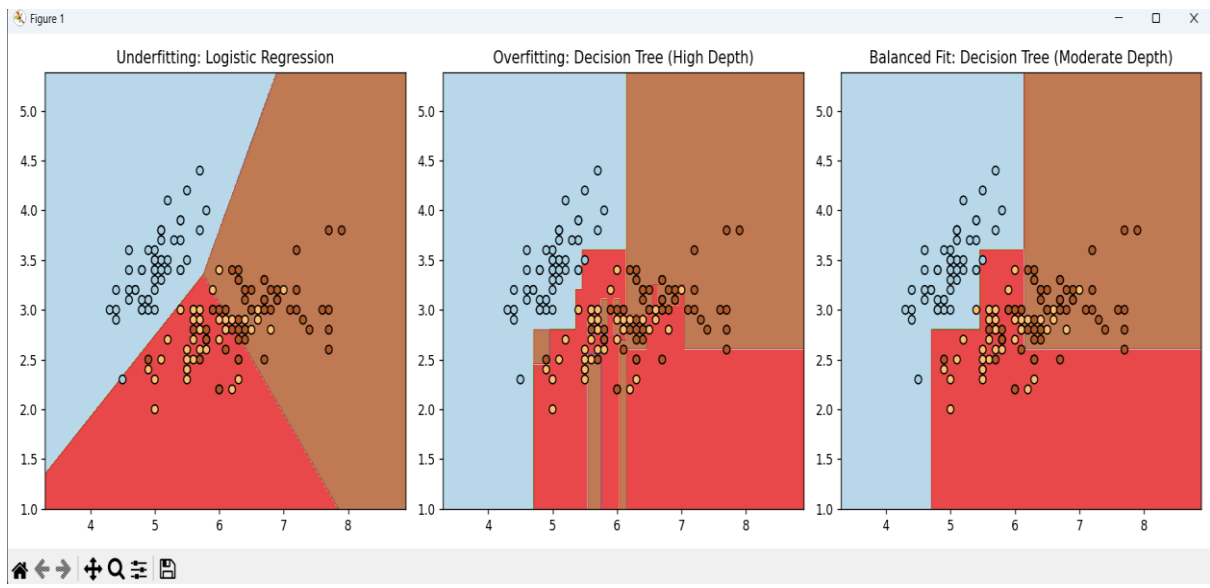
## Output:

```

PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS  SQL CONSOLE
[Running] python -u "c:\Users\harsh\OneDrive\Desktop\PR Project\first1.1..py"
Dataset saved as 'iris_dataset.csv'
Underfitting Model: Train Accuracy = 0.7904761904761904 , Test Accuracy = 0.8222222222222222
Overfitting Model: Train Accuracy = 0.9428571428571428 , Test Accuracy = 0.7111111111111111
Balanced Model: Train Accuracy = 0.8285714285714286 , Test Accuracy = 0.7555555555555555

[Done] exited with code=0 in 72.852 seconds

```



## [Question 1.2]

### Implementation:

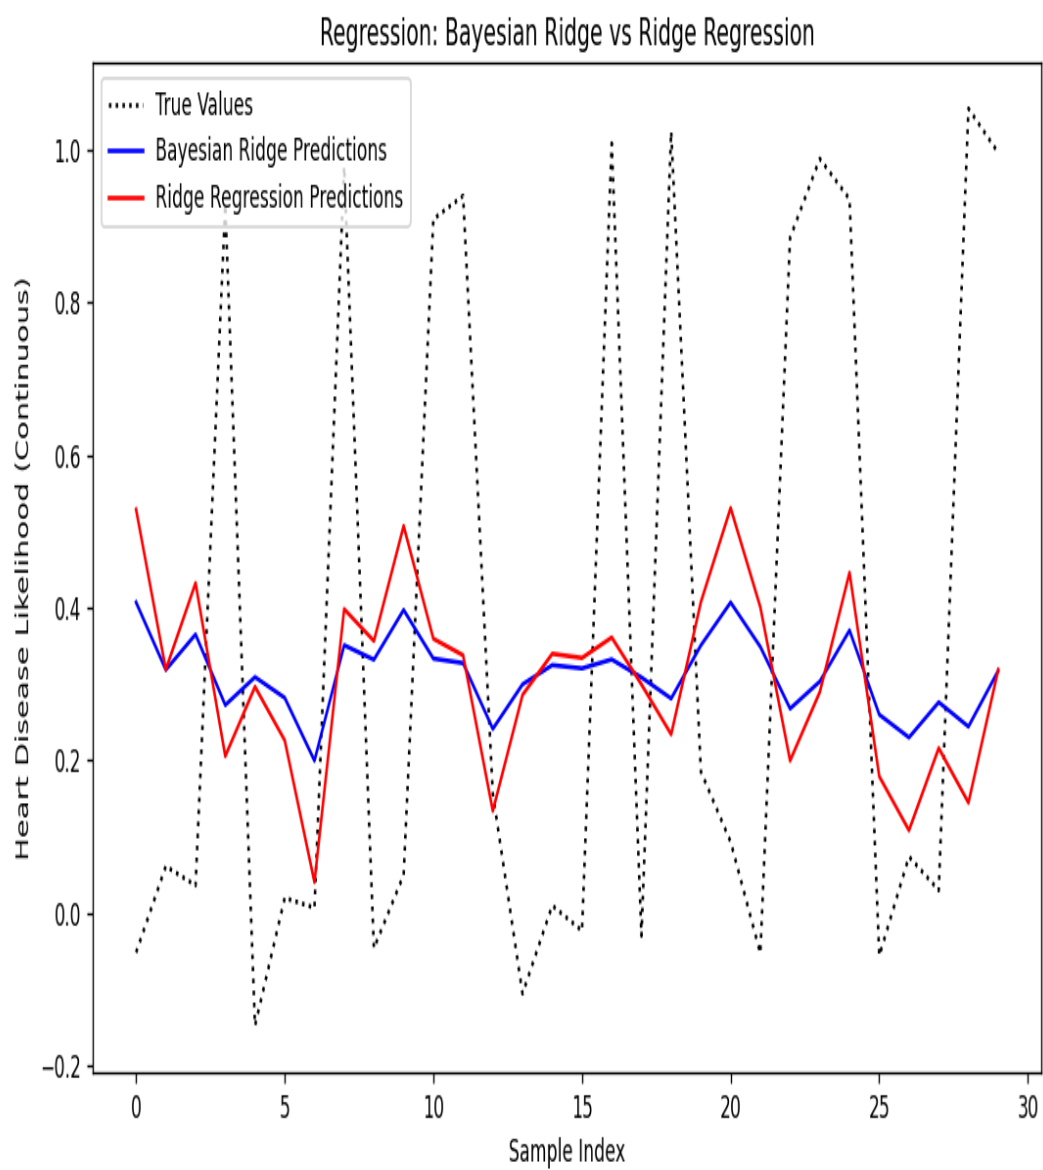
```
second1.2.py > ...
1  import pandas as pd
2  import numpy as np
3  from sklearn.model_selection import train_test_split
4  from sklearn.preprocessing import StandardScaler
5  from sklearn.linear_model import BayesianRidge, Ridge
6  from sklearn.metrics import mean_squared_error, r2_score
7  import matplotlib.pyplot as plt
8
9  # Load the synthetic heart disease dataset
10 dataset_path = "synthetic_heart_disease_dataset.csv"
11 data = pd.read_csv(dataset_path)
12
13 # Modify the target to be continuous (for regression)
14 np.random.seed(42)
15 data["HeartDiseaseLikelihood"] = data["HeartDisease"] + np.random.normal(0, 0.1, len(data))
16
17 # Features and target
18 X = data[["Age", "Cholesterol", "MaxHeartRate"]]
19 y = data["HeartDiseaseLikelihood"]
20
21 # Split the dataset
22 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
23
24 # Standardize the data
25 scaler = StandardScaler()
26 X_train = scaler.fit_transform(X_train)
27 X_test = scaler.transform(X_test)
28
29 # Bayesian Ridge Regression
30 bayesian_model = BayesianRidge()
31 bayesian_model.fit(X_train, y_train)
32 y_pred_bayesian = bayesian_model.predict(X_test)
33
34 # Ridge Regression (LS Regularization)
35 ridge_model = Ridge(alpha=1.0)
36 ridge_model.fit(X_train, y_train)
37 y_pred_ridge = ridge_model.predict(X_test)
38
39 # Evaluation Metrics
40 mse_bayesian = mean_squared_error(y_test, y_pred_bayesian)
41 r2_bayesian = r2_score(y_test, y_pred_bayesian)
42
43 mse_ridge = mean_squared_error(y_test, y_pred_ridge)
44 r2_ridge = r2_score(y_test, y_pred_ridge)
45
46
47 # Evaluation Metrics
48 mse_bayesian = mean_squared_error(y_test, y_pred_bayesian)
49 r2_bayesian = r2_score(y_test, y_pred_bayesian)
50
51 mse_ridge = mean_squared_error(y_test, y_pred_ridge)
52 r2_ridge = r2_score(y_test, y_pred_ridge)
53
54 # Display Metrics
55 print(f"Bayesian Ridge Regression - MSE: {mse_bayesian:.4f}, R2 Score: {r2_bayesian:.4f}")
56 print(f"Ridge Regression (LS Regularization) - MSE: {mse_ridge:.4f}, R2 Score: {r2_ridge:.4f}")
57
58 # Visualization
59 plt.figure(figsize=(10, 6))
60
61 # True Values vs Predictions
62 plt.plot(y_test.values, label="True Values", color="black", linestyle="dotted")
63 plt.plot(y_pred_bayesian, label="Bayesian Ridge Predictions", color="blue")
64 plt.plot(y_pred_ridge, label="Ridge Regression Predictions", color="red")
65
66 plt.legend()
67 plt.title("Regression: Bayesian Ridge vs Ridge Regression")
68 plt.xlabel("Sample Index")
69 plt.ylabel("Heart Disease Likelihood (Continuous)")
70 plt.show()
71
```

## Output:

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS  SQL CONSOLE

[Running] python -u "c:\Users\harsh\OneDrive\Desktop\PR Project\second1.2.py"
Bayesian Ridge Regression - MSE: 0.2255, R2 Score: -0.0357
Ridge Regression (LS Regularization) - MSE: 0.2421, R2 Score: -0.1115
```

Figure 1



---

### [Question 1.3]

#### Solution:

### Question 1: Classification with the Iris Dataset

**Dataset:** The Iris dataset contains 3 species of iris flowers (Setosa, Versicolor, Virginica) and 4 features (sepal length, sepal width, petal length, petal width). This dataset is used to demonstrate underfitting, overfitting, and the bias-variance tradeoff.

#### **Models:**

1. **Underfitting:** Logistic Regression (low complexity, linear decision boundaries)
  - **Result:** Struggles to separate species, leading to low accuracy on both training and testing sets.
2. **Overfitting:** Decision Tree with high depth (no limit on tree depth)
  - **Result:** Memorizes training data (overfitting), high accuracy on training but poor test accuracy.
3. **Balanced Fit:** Decision Tree with moderate depth (e.g., depth=3)
  - **Result:** Generalizes well, balancing accuracy on both training and testing sets.

#### **Bias-Variance Tradeoff:**

- **Underfitting:** High bias (simple model, misses complex patterns)
- **Overfitting:** High variance (complex model, sensitive to noise in training data)
- **Balanced Model:** Low bias and low variance, performs well on both training and testing sets.

#### **Visualizations:**

- **Decision Boundaries:** Logistic Regression creates simple straight lines, while the high-depth decision tree has jagged, overfit boundaries. The balanced decision tree creates smooth, generalizable boundaries.
- **Train vs Test Accuracy:** Underfitting shows low accuracy for both, overfitting shows a high training accuracy and low test accuracy, and the balanced model performs well for both.

---

### Question 2: Regression with the Heart Disease Dataset

**Dataset:** The Heart Disease dataset is used to predict the likelihood of heart disease based on features such as age, cholesterol, and max heart rate. We'll treat this as a regression problem (predicting a probability between 0 and 1).

#### **Models:**

1. **Bayesian Regression:** Incorporates uncertainty in predictions, providing a probabilistic distribution over possible outcomes.
  - **Result:** Offers robust predictions, useful for understanding uncertainty in predictions.
2. **Ridge Regression (LS Regularization):** Applies L2 regularization to penalize large coefficients and prevent overfitting.
  - **Result:** Regularization improves generalization, leading to better performance on unseen data.

### Preprocessing:

- Features are standardized to ensure consistent scaling, which is important for Ridge regression.

### Impact of LS Regularization:

- **Ridge Regression:** Reduces the influence of irrelevant features and prevents overfitting.
- **Bayesian Regression:** Uses prior distributions to handle overfitting and manage uncertainty.

### Evaluation Metrics:

- **MSE (Mean Squared Error):** Both Ridge and Bayesian regressions reduce MSE compared to unregularized models, showing better prediction accuracy.
- **R<sup>2</sup> Score:** Regularized models show higher R<sup>2</sup>, meaning they explain more of the variance in the data.

### Visualizations:

- **True vs Predicted Values:** Regularized models provide smoother, more accurate predictions. Unregularized models may show erratic predictions, indicating overfitting.
- **Feature Importance:** Ridge regression shows a reduced importance for irrelevant features, resulting in a simpler, more generalizable model.

---

### [Question 2.1]

### Implementation:

### SVM for Classification

fourth2.1classification.py > ...

```
1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn import datasets
4  from sklearn.model_selection import train_test_split
5  from sklearn.svm import SVC
6  from sklearn.metrics import accuracy_score
7
8  # Load Iris dataset
9  data = datasets.load_iris()
10 X = data.data[:, :2] # Using only the first two features for simplicity
11 y = data.target
12
13 # Split dataset into train and test
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
15
16 # Linear SVM
17 linear_svm = SVC(kernel='linear')
18 linear_svm.fit(X_train, y_train)
19 y_pred_linear = linear_svm.predict(X_test)
20
21 # Non-Linear SVM with RBF kernel
22 rbf_svm = SVC(kernel='rbf')
23 rbf_svm.fit(X_train, y_train)
24 y_pred_rbf = rbf_svm.predict(X_test)
25
26 # Evaluate accuracy
27 print("Linear SVM Accuracy: ", accuracy_score(y_test, y_pred_linear))
28 print("RBF SVM Accuracy: ", accuracy_score(y_test, y_pred_rbf))
29
30 # Plot decision boundaries
31 def plot_decision_boundary(X, y, model, title):
32     h = .02 # Step size in the mesh
33     x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
34     y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
35     xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
36                          np.arange(y_min, y_max, h))
37     Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
38     Z = Z.reshape(xx.shape)
39
40     plt.contourf(xx, yy, Z, alpha=0.75, cmap=plt.cm.coolwarm)
41     plt.scatter(X[:, 0], X[:, 1], c=y, s=30, edgecolor='k', cmap=plt.cm.coolwarm)
42     plt.title(title)
43     plt.show()
44
45 # Plot decision boundaries for Linear and RBF SVM
46 plt.figure(figsize=(12, 6))
47
48 plt.subplot(1, 2, 1)
49 plot_decision_boundary(X, y, linear_svm, "Linear SVM Decision Boundary")
50
51 plt.subplot(1, 2, 2)
52 plot_decision_boundary(X, y, rbf_svm, "RBF SVM Decision Boundary")
```



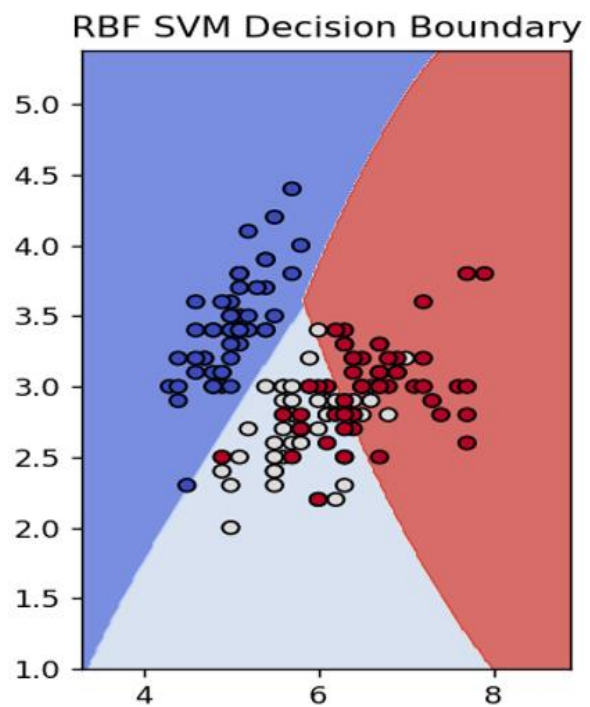
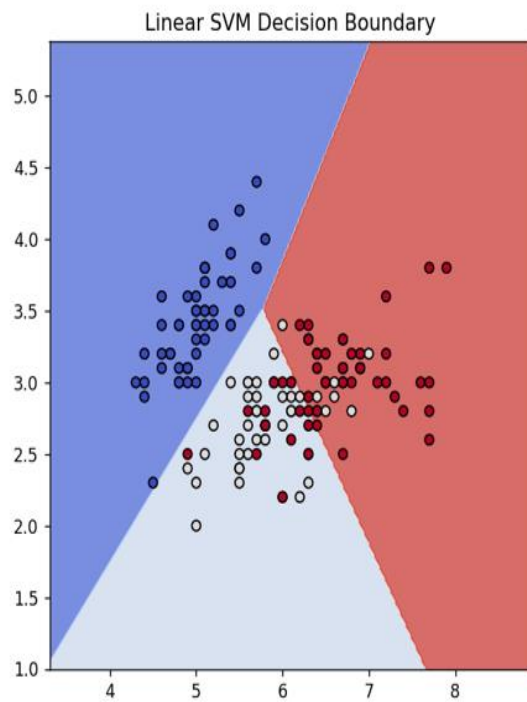
## SVM for Regression

```
fourth2.1regression.py > ...
1  from sklearn.svm import SVR
2  from sklearn.datasets import load_diabetes
3  from sklearn.preprocessing import StandardScaler
4  import matplotlib.pyplot as plt
5
6  # Load the Diabetes dataset (a regression dataset)
7  diabetes = load_diabetes()
8  X, y = diabetes.data[:, 2:3], diabetes.target # Using a single feature for easy visualization
9
10 # Standardize the features
11 scaler = StandardScaler()
12 X_scaled = scaler.fit_transform(X)
13
14 # Linear SVR
15 linear_svr = SVR(kernel='linear')
16 linear_svr.fit(X_scaled, y)
17
18 # Non-Linear SVR with RBF kernel
19 rbf_svr = SVR(kernel='rbf')
20 rbf_svr.fit(X_scaled, y)
21
22 # Plotting the results
23 plt.figure(figsize=(12, 6))
24
25 # Plot for Linear SVR
26 plt.subplot(1, 2, 1)
27 plt.scatter(X_scaled, y, color='gray', label='Data')
28 plt.plot(X_scaled, linear_svr.predict(X_scaled), color='red', label='Linear SVR')
29 plt.title("Linear SVM Regression")
30 plt.legend()
31
32 # Plot for RBF SVR
33 plt.subplot(1, 2, 2)
34 plt.scatter(X_scaled, y, color='gray', label='Data')
35 plt.plot(X_scaled, rbf_svr.predict(X_scaled), color='blue', label='RBF SVR')
36 plt.title("RBF SVM Regression")
37 plt.legend()
38
39 plt.show()
40
```

## Output:

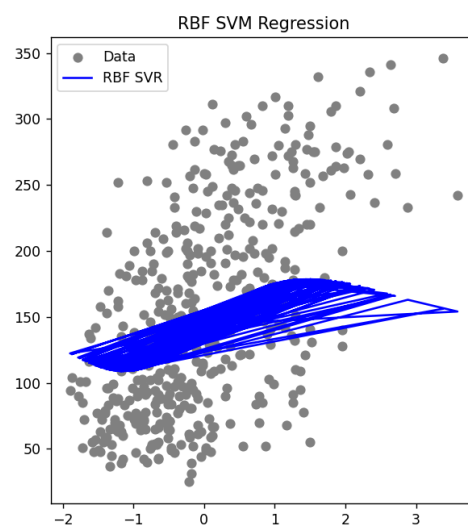
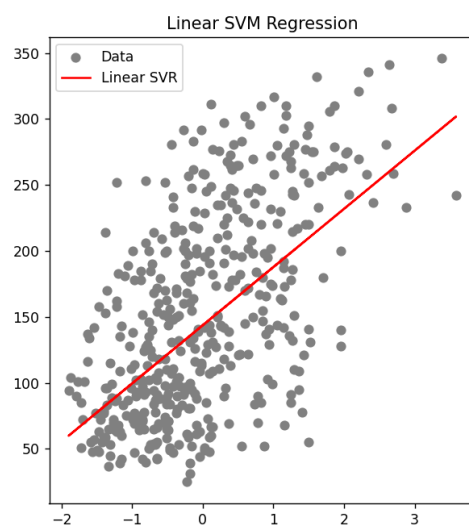
## SVM for Classification

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS  SQL CONSOLE
[Running] python -u "c:\Users\harsh\OneDrive\Desktop\PR Project\fourth2.1classification.py"
Linear SVM Accuracy: 0.8
RBF SVM Accuracy: 0.8
```



## SVM for Regression

Figure 1



---

## **Illustration**

### **[1.] Classification using SVM with the Iris Dataset**

#### **Dataset:**

The **Iris dataset** is a popular dataset in machine learning, used to classify iris flowers into three species (Setosa, Versicolor, and Virginica) based on four features: sepal length, sepal width, petal length, and petal width.

#### **Steps:**

1. **Load the Iris dataset** and split it into training and testing sets.
2. **Train SVM models** with both **linear kernel** and **non-linear kernel (RBF)**.
3. **Evaluate** the models using accuracy and visualize their decision boundaries.

#### **Key Observations:**

- **Linear SVM:** The decision boundary will be a straight line (or hyperplane in higher dimensions). Since SVM is trying to maximize the margin between classes, it will work well if the classes are linearly separable.
- **RBF SVM:** The decision boundary will be more flexible, allowing for a non-linear separation. It works well when the classes are not linearly separable and uses kernel tricks to map the data into a higher-dimensional space.

---

### **[2.] Regression using SVM with the Diabetes Dataset**

#### **Dataset:**

In this example, we will use the **Diabetes dataset** and apply **Support Vector Regression (SVR)** in two different modes: **Linear** and **RBF (Radial Basis Function)** kernels. Our goal is to predict **disease progression** based on a single feature (like BMI or age) for easy visualization.

#### **Steps:**

1. **Load and prepare the data:**
  - The dataset contains multiple features, but for simplicity, we'll use just one feature ( $X[:,2:3]$  for easy visualization).
2. **Standardize the features:**
  - The Support Vector Machines (SVM) are sensitive to the scale of data, so standardization is necessary to ensure all features have the same scale.
3. **Apply SVR models:**
  - **Linear SVR:** Assumes a linear relationship between the input and target.

- **RBF SVR**: A non-linear kernel that can capture complex relationships between the data.
- 4. **Visualize the predictions**:
  - Show the **data points** and compare the predictions made by both models.

## Key Observations:

1. **Linear SVR** might perform poorly on complex datasets where the relationship between features and target is not linear.
2. **RBF SVR** is more powerful in capturing non-linear relationships, leading to better performance in cases like this where data complexity requires flexibility in the model.

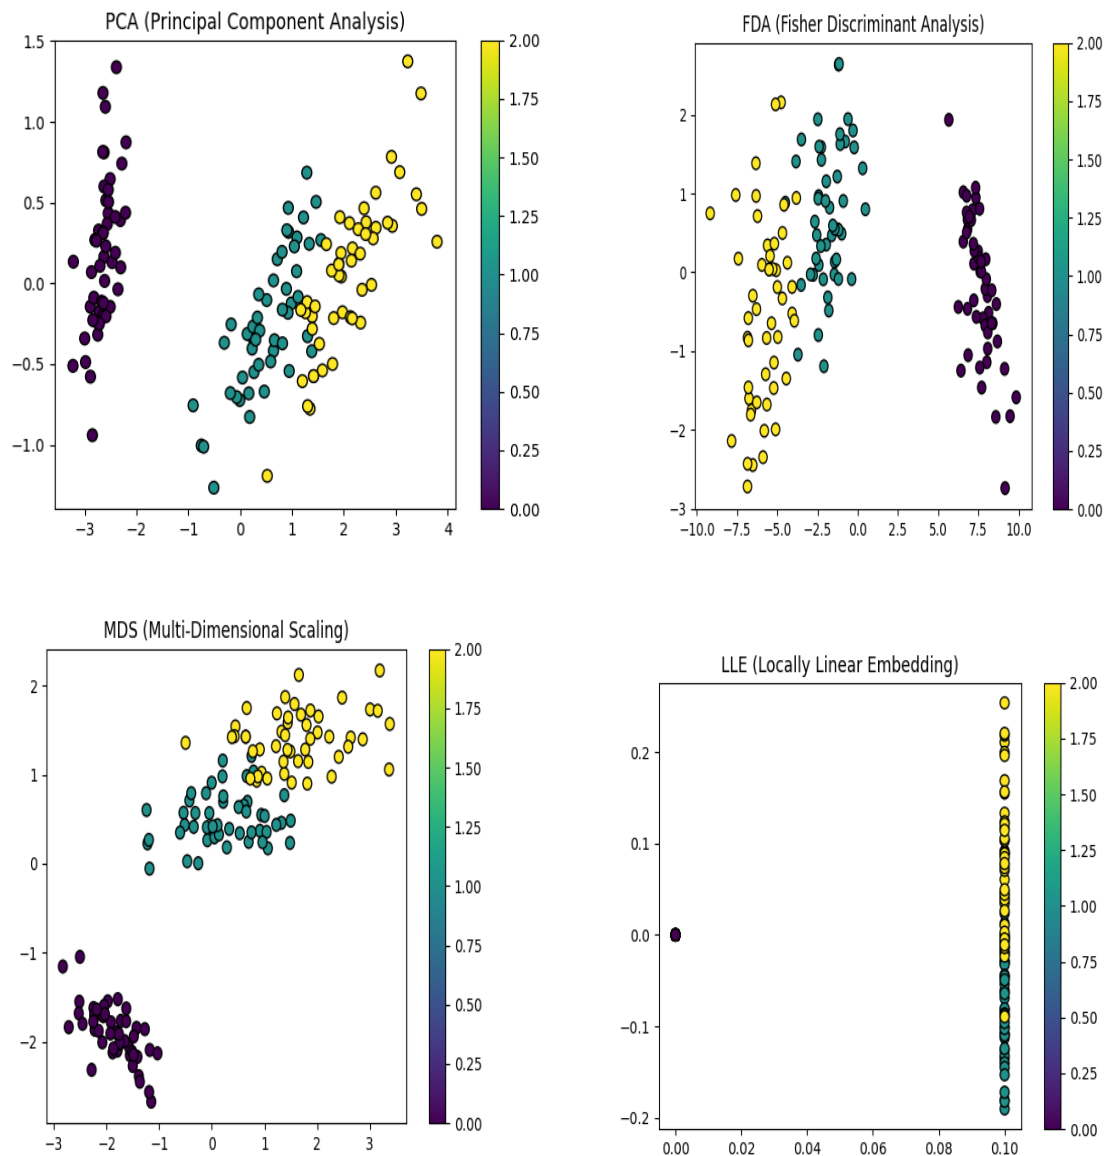
---

## [Question 2.2]

**Implementation:** we apply PCA, FDA, MDA, LLE on the **Iris Dataset**, which is often used for classification tasks and perform dimensionality reduction to two dimensions, followed by visualization and evaluation.

```
fifth2.2.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn import datasets
4  from sklearn.decomposition import PCA
5  from sklearn.preprocessing import LabelEncoder
6  from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
7  from sklearn.manifold import MDS
8  from sklearn.neighbors import NearestNeighbors
9  from sklearn.metrics import pairwise_distances
10
11 # Load the Iris dataset
12 iris = datasets.load_iris()
13 X = iris.data
14 y = iris.target
15
16 # Create a scatter plot function for comparing results
17 def scatter_plot(X_transformed, title):
18     plt.scatter(X_transformed[:, 0], X_transformed[:, 1], c=y, cmap='viridis', edgecolor='k', s=50)
19     plt.title(title)
20     plt.colorbar()
21     plt.show()
22
23 # 1. PCA - Principal Component Analysis
24 pca = PCA(n_components=2)
25 X_pca = pca.fit_transform(X)
26 scatter_plot(X_pca, "PCA (Principal Component Analysis)")
27
28 # 2. FDA - Fisher Discriminant Analysis
29 lda = LDA(n_components=2)
30 X_lda = lda.fit_transform(X, y)
31 scatter_plot(X_lda, "FDA (Fisher Discriminant Analysis)")
32
33 # 3. MDA - Multi-Dimensional Scaling (MDS)
34 mds = MDS(n_components=2, random_state=42)
35 X_mds = mds.fit_transform(X)
36 scatter_plot(X_mds, "MDS (Multi-Dimensional Scaling)")
37
38 # 4. LLE - Locally Linear Embedding
39 from sklearn.manifold import LocallyLinearEmbedding
40
41 lle = LocallyLinearEmbedding(n_components=2, n_neighbors=10)
42 X_lle = lle.fit_transform(X)
43 scatter_plot(X_lle, "LLE (Locally Linear Embedding)")
```

## Output:



## 1. Principal Component Analysis (PCA)

- **Purpose:** PCA is an unsupervised technique that reduces the dimensionality by finding the directions (principal components) where the data has the most variance.
- **Key Idea:** PCA looks for linear combinations of the original features that capture the most variance in the data.
- **Advantages:** Simple and effective for reducing dimensions in high-dimensional data.
- **Limitations:** Does not consider class labels, so may not be optimal for classification tasks.

## 2. Fisher Discriminant Analysis (FDA)

- **Purpose:** FDA is a supervised technique that reduces the dimensionality by maximizing the separation between classes.
- **Key Idea:** FDA tries to find a lower-dimensional space that best separates the classes in the dataset.
- **Advantages:** Works well for classification problems, as it takes class labels into account.
- **Limitations:** Assumes that the data from each class is normally distributed with the same covariance matrix, which may not always be true.

## 3. Multi-Dimensional Scaling (MDA)

- **Purpose:** MDA is an unsupervised technique that tries to preserve the pairwise distances between data points in a lower-dimensional space.
- **Key Idea:** MDA minimizes the stress function, which is a measure of how well the distances are preserved in the lower-dimensional space.
- **Advantages:** Useful when the data is represented by a dissimilarity matrix.
- **Limitations:** May not scale well to large datasets and is sensitive to the initial configuration.

## 4. Locally Linear Embedding (LLE)

- **Purpose:** LLE is a nonlinear dimensionality reduction technique that focuses on preserving local structures in the data.
- **Key Idea:** LLE works by finding a lower-dimensional representation where each point's neighbors are preserved as linear combinations.
- **Advantages:** Effective for reducing the dimensionality of nonlinear datasets.
- **Limitations:** Sensitive to noise and can be computationally expensive.

---

### [Question 2.3]

#### Solution:

#### [1.] SVM Implementation and Illustration

- **Support Vector Machines (SVM):**
  - Implemented SVM for both classification and regression tasks using different kernels: **Linear** and **RBF**.
  - For **classification**, SVM with the Iris dataset demonstrated how linear and non-linear decision boundaries separate data points belonging to different classes.
  - For **regression**, SVM with the Diabetes dataset illustrated how models predict continuous outputs, with linear and RBF kernels adapting to different data distributions.
- **Key Takeaways:**
  - The **linear kernel** works well for linearly separable data but struggles with non-linear patterns.

- The **RBF kernel** effectively handles non-linear relationships by mapping data into higher-dimensional spaces, though it may require hyperparameter tuning for optimal performance.
- 

## [2.] Dimensionality Reduction Comparative Analysis

- **Techniques Evaluated:** PCA, FDA, MDA, and LLE on the Iris dataset.
  - **Observations:**
    - **PCA (Principal Component Analysis):**
      - Unsupervised method; reduces dimensions by preserving maximum variance.
      - Effective for linear datasets but does not leverage class labels, which can affect classification performance.
    - **FDA (Fisher Discriminant Analysis):**
      - Supervised method; maximizes class separability in reduced dimensions.
      - Ideal for classification tasks but assumes normality in class distributions.
    - **MDA (Multi-Dimensional Scaling):**
      - Unsupervised; preserves pairwise distances between data points.
      - Useful for visualizing relationships but less effective for classification.
    - **LLE (Locally Linear Embedding):**
      - Non-linear, unsupervised method; preserves local structures.
      - Captures complex patterns in non-linear datasets but can be computationally expensive.
  - **Key Takeaways:**
    - For **linear relationships**, PCA and FDA perform well, with FDA being superior for classification.
    - For **non-linear relationships**, LLE outperforms other techniques by preserving local structures.
    - MDA is useful for exploratory data analysis where distance preservation is critical.
-